ILS - Z 604: Assignment #3

Due on Tuesday, April 19, 2016

Prof. Xiazhong Liu

vpatani (Vivek Patani)

Contents

Task 1

Listing 1 Shows the IPynb Script.

Listing 1: Shows the IPynb Script for cleaning

```
#!/usr/bin/perl

#Part 1
import os
import graphlab
graphlab.product_key.set_product_key('Your API Key')
location = 'http://www.vivekpatani.tk/resources/reuters.csv'
sf = graphlab.SFrame.read_csv(location, header=False)
```

- How to configure the graphLab for AWS:
  - Register as a student and Dato will send you a key.
  - Register the key as *graphlab.product − key* function as shown above.
  - Execute in order to register

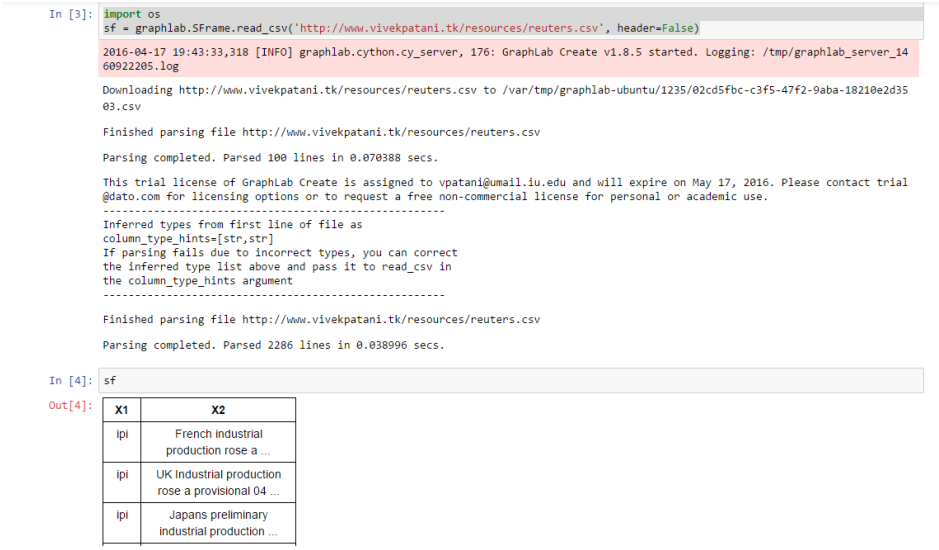

```
In [3]: import os
sf = graphlab.SFrame.read_csv('http://www.vivekpatani.tk/resources/reuters.csv', header=False)

2016-04-17 19:43:33,318 [INFO] graphlab.cython.cy_server, 176: GraphLab Create v1.8.5 started. Logging: /tmp/graphlab_server_14
60922205.log

Downloading http://www.vivekpatani.tk/resources/reuters.csv to /var/tmp/graphlab-ubuntu/1235/02cd5fbc-c3f5-47f2-9aba-18210e2d35
03.csv

Finished parsing file http://www.vivekpatani.tk/resources/reuters.csv

Parsing completed. Parsed 100 lines in 0.070388 secs.

This trial license of GraphLab Create is assigned to vpatani@umail.iu.edu and will expire on May 17, 2016. Please contact trial
@dato.com for licensing options or to request a free non-commercial license for personal or academic use.
------------------------------------------------------
Inferred types from first line of file as
column_type_hints=[str,str]
If parsing fails due to incorrect types, you can correct
the inferred type list above and pass it to read_csv in
the column_type_hints argument
------------------------------------------------------

Finished parsing file http://www.vivekpatani.tk/resources/reuters.csv

Parsing completed. Parsed 2286 lines in 0.038996 secs.

In [4]: sf
```

Out[4]:

X1	X2
ipi	French industrial production rose a ...
ipi	UK industrial production rose a provisional 04 ...
ipi	Japans preliminary industrial production ...

Listing 2 Shows the IPynb Script.

Listing 2: Shows the IPynb Script for counting

```
#Part 2
word_list = graphlab.text_analytics.count_words(sf['X2'])
docs = sf['word_list'].dict_trim_by_values(2)
docs = docs.dict_trim_by_keys(graphlab.text_analytics.stopwords(),exclude=True)
```

- Now we count the words and eliminate the stop words and words not crossing threshold
  - Use the *SFrame.read − csv* function giving input of the location of your file.
  - View sf to confirm the number of rows.

Task 1 continued on next page...

– Once we have the data ready, we use the libraries to count the words that exist in the dataset. We use an Sarray.SArray to store the word and their count.

– The next step is to see what words have appeared lesser than 2 times and we pass a parameter as 2 to the function $trim - by - value$.

– Next step is to eliminate the stop word. We pass *doc* as the input and the the output is over ridden on *doc* itself by passing it through *stopwords* function.

– Just to make sure we've got it right, just check *doc*[0].

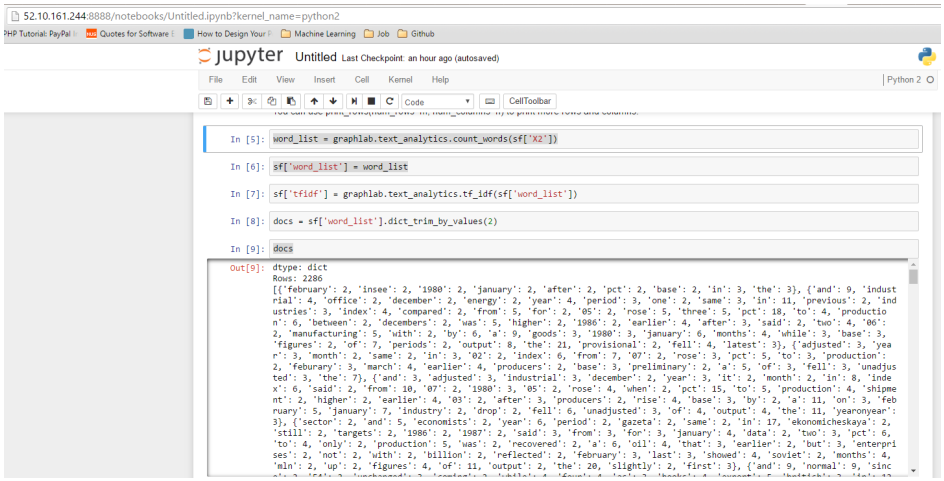

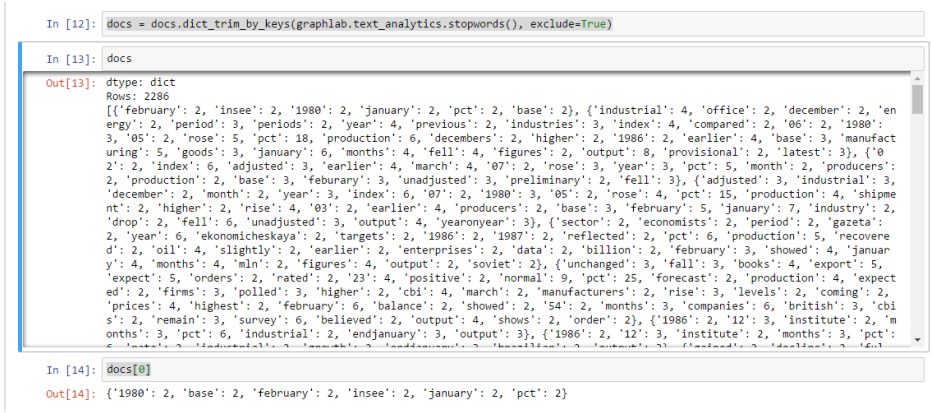

- This completes the first part of **cleaning data** and **generating features**.

Task 2

Listing 3 Shows the IPynb Script.

Listing 3: Shows the IPynb Script for generating Features

```
#Part 3
word_list = graphlab.text_analytics.count_words(sf['X2'])
sf['word_list'] = word_list
sf['tfidf'] = graphlab.text_analytics.tf_idf(sf['word_list'])
```

- Here the we simply create a list(bag) of words for each document.
- This is straightforward since we just generate a word list by the $count_words$ function by giving input as the the feature column over which we need to do this.
- Next step is to add it the frame generated by us.
- Next run the TF-IDF from the analytics library of graph lab over the word list generated.
- We generated two things
  1. Word List
  2. TF-IDF Values
- This is how the Frame will look like

In [11]: sf

Out[11]:

X1	X2	word_list	tfidf
ipi	French industrial production rose a ...	{'and': 1, 'adjusted': 1, 'industrial': 1, '198': ...	{'and': 0.13967445713823656, ...
ipi	UK Industrial production rose a provisional 04 ...	{'six': 1, 'office': 2, 'show': 1, '1264': 1, ...	{'six': 2.5034502275001693, ...
ipi	Japans preliminary industrial production ...	{'and': 1, 'adjusted': 3, 'industrial': 1, 'goo ...	{'and': 0.13967445713823656, ...
ipi	Japans industrial production index base ...	{'adjusted': 3, 'trade': 1, 'previous': 1, ...	{'adjusted': 8.917154728671, 'trade': ...
ipi	The Soviet economy recovered slightly last ...	{'sector': 2, 'all': 1, 'show': 1, 'obligatio ...	{'sector': 6.1803358904267665, ...
ipi	British manufacturers expect output to grow ...	{'rating': 1, 'all': 1, 'reuter': 1, 'consider': ...	{'rating': 6.348264483234865, 'a ...
ipi	Industrial output in January was 609 pct a ...	{'and': 2, 'reuter': 1, 'show': 1, 'trend': 1, ...	{'and': 0.2793489142764731, ...
ipi	The growth rate of Brazilian industrial ...	{'and': 2, 'industrial': 2, 'show': 1, 'trend' ...	{'and': 0.2793489142764731, ...
ipi	Swiss industrial output rose nine pct in the ...	{'and': 4, 'industrial': 1, 'stood': 1, 'office': ...	{'and': 0.5586978285529463, ...
ipi	Chinas industrial output rose 141 pct in the f ...	{'and': 1, 'seven': 1, 'industrial': 3, '1987': ...	{'and': 0.13967445713823656, ...

- Now, we have attached the **features** that will help us analyse the data.

Task 3

Listing 4 Shows the IPynb Script.

Listing 4: Shows the IPynb Script for generating Model & Evaluating it

```
#Part 4
train_data, test_data = sf.random_split(0.8)
test_data
len(sf)
len(test_data)
len(train_data)
```

- In this we have cleaned the data and build the elements required to build a model.
- We first divide the set into **Train** & **Test** in order to build and evaluate the model. We split the data in the ratio of 80:20, where 80 is train and 20 is test.
- This is how the Train & Test Frame will look like

```
In [17]: train_data, test_data = sf.random_split(0.8)

In [19]: test_data
```

Out[19]:

X1	X2	word_list	tfidf
ipi	Swiss industrial output rose nine pct in the ...	{'and': 4, 'industrial': 1, 'stood': 1, 'office': ...	{'and': 0.5586978285529463, ...
ipi	Swedish industrial production rose 26 pc ...	{'and': 2, 'spell': 1, 'all': 1, 'industrial': ...	{'and': 0.2793489142764731, ...
ipi	US industrial production rose 05 pct in February ...	{'six': 1, 'month': 3, 'paper': 1, 'still': 1, ...	{'six': 2.5034502275001693, ...
ipi	The Bank of France expects a continued ...	{'sector': 4, 'all': 2, 'unemployment': 1, ...	{'sector': 12.360671780853533, ...
ipi	Canadas industrial product price index rose ...	{'yearly': 1, 'and': 3, 'industrial': 1, ...	{'yearly': 5.169609486893219, 'a ...
ipi	British manufacturers expect output to grow ...	{'rating': 1, 'all': 1, 'consider': 1, 'polled': ...	{'rating': 6.348264483234865, 'a ...
ipi	Industrial production rose 48 pct on a ...	{'and': 1, 'industrial': 3, 'show': 1, 'decemb ...	{'and': 0.13967445713823656, ...
ipi	Japans preliminary industrial production ...	{'and': 1, 'adjusted': 3, 'industrial': 1, 'goo ...	{'and': 0.13967445713823656, ...
ipi	Swedish industrial production rose 15 pc ...	{'sector': 1, 'and': 2, 'industrial': 1, ...	{'sector': 3.0901679452133832, ...
ipi	A leading Soviet economist said the ...	{'years': 1, 'inflating': 1, 'economist': 1, ...	{'years': 1.8624410548793402, ...

[? rows x 4 columns]
Note: Only the head of the SFrame is printed. This SFrame is lazily evaluated.
You can use len(sf) to force materialization.

- The length of each is

```
In [20]: len(sf)
Out[20]: 2286

In [21]: len(test_data)
Out[21]: 450

In [22]: len(train_data)
Out[22]: 1836
```

- We use train to learn the model and data.

Task 3 continued on next page. . .

- Test will be used to evaluate features of the model such as accuracy.
- Now we start building models with different approaches:
  1. Boosted Tree Classifier: Listing 5 Shows the IPynb Script.

Listing 5: Shows the IPynb Script for generating Model & Evaluating it

```
model = graphlab.boosted_trees_classifier.create(train_data, target='X1',
                                                 max_iterations=2,
                                                 max_depth = 3)
predictions = model.classify(test_data)
results = model.evaluate(test_data)
```

Task 3 continued on next page...

– After training the data with the given model.

– Output:

```
Boosted trees classifier:

--------------------------------------------------------

Number of examples          : 1756

Number of classes           : 34

Number of feature columns   : 1

Number of unpacked features : 17269

Create disk column page 1/1

+-----------+--------------+-------------------+-------------------+---------------------+---------------------+
| Iteration | Elapsed Time | Training-accuracy | Training-log_loss | Validation-accuracy | Validation-log_loss |
+-----------+--------------+-------------------+-------------------+---------------------+---------------------+
| 1         | 10.136764    | 0.876993          | 1.235948          | 0.575000            | 2.121445            |
| 2         | 19.866346    | 0.890091          | 0.934086          | 0.525000            | 1.927115            |
| 3         | 29.492107    | 0.899772          | 0.741076          | 0.537500            | 1.795780            |
| 4         | 39.086573    | 0.910592          | 0.603383          | 0.512500            | 1.716426            |
| 5         | 48.789179    | 0.917426          | 0.499902          | 0.512500            | 1.668092            |
| 6         | 58.514369    | 0.919134          | 0.421551          | 0.525000            | 1.626862            |
| 7         | 68.342350    | 0.921982          | 0.361099          | 0.525000            | 1.573787            |
| 8         | 78.620567    | 0.922551          | 0.314064          | 0.537500            | 1.547404            |
| 9         | 89.087766    | 0.923690          | 0.275237          | 0.537500            | 1.539655            |
| 10        | 99.270443    | 0.924260          | 0.245745          | 0.537500            | 1.526494            |
+-----------+--------------+-------------------+-------------------+---------------------+---------------------+

PROGRESS: Creating a validation set from 5 percent of training data. This may take a while.
          You can set ``validation_set=None`` to disable validation tracking.
```

– After running the test, the accuracy is 71%:

– Output:

```
In [12]: results = model.evaluate(test_data)

         External memory mode: 1 batches

In [13]: results

Out[13]: {'accuracy': 0.7103004291845494,
          'auc': 0.9698714070751959,
          'confusion_matrix': Columns:
                 target_label    str
                 predicted_label str
                 count   int

          Rows: 88

          Data:
          +----------------+------------------+-------+
          |  target_label  | predicted_label  | count |
          +----------------+------------------+-------+
          |    pet-chem    |     pet-chem     |   3   |
          |    veg-oil     |  oilseedsoybean  |   4   |
          |   graincorn    |    graincorn     |   8   |
          |      jobs      |       gnp        |   2   |
          |      jobs      |       jobs       |   5   |
          |  dlrmoney-fx   | interestmoney-fx |   1   |
          |    tradebop    |     tradebop     |   2   |
          |      gnp       |       cpi        |   2   |
          |    reserves    |       gold       |   2   |
          | money-fxdlryen |   money-fxdlr    |   6   |
          +----------------+------------------+-------+
          [88 rows x 3 columns]
          Note: Only the head of the SFrame is printed.
          You can use print_rows(num_rows=m, num_columns=n) to print more rows and columns.,
          'f1_score': 0.6662522123798048,
          'log_loss': 1.4751669587448086,
          'precision': 0.7308917606173222,
          'recall': 0.6861378490790255,
          'roc_curve': Columns:
                 threshold       float
                 fpr     float
                 tpr     float
                 p       int
                 n       int
                 class   int

          Rows: 3400034
```

Task 3 continued on next page...

– The prediction excerpt
– Output:

```
In [16]: predictions
```

Out[16]:

class	probability
ipi	0.463541865349
ipi	0.495206207037
ipi	0.49030277133
ipi	0.490942507982
ipi	0.494604974985
ipi	0.397208571434
ipi	0.0526305325329
alum	0.338211506605
ipi	0.490942507982
reserves	0.563541531563

[466 rows x 2 columns]
Note: Only the head of the SFrame is p
You can use print_rows(num_rows=m,